kafka学习笔记-learning1

环境安装（Linux上的安装使用）

环境及软件版本

- LINUX : CentSO-8.2.2004-x86_64
- JDK : 11.0.8
- Zookeeper : 3.6.2
  下载地址：https://mirror.bit.edu.cn/apache/zookeeper/zookeeper-3.6.2/apache-zookeeper-3.6.2-bin.tar.gz
- kafka : 2.13-2.6.0
  下载地址：https://mirror.bit.edu.cn/apache/kafka/2.6.0/kafka_2.12-2.5.0.tgz

软件安装

- Java 安装

  ```
  yum install java-11-openjdk-devel.x86_64
  ```

  检查安装情况

  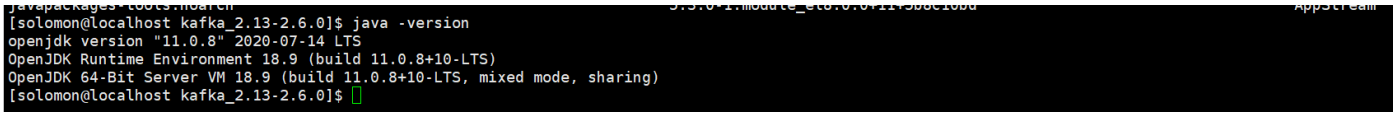

- 安装zookeeper
  到安装包目录下执行命令

  ```
  tar -zxvf apache-zookeeper-3.6.2-bin.tar.gz
  ```

  zookeeper 配置

  ```
  cp zoo_sample.cfg zoo.cfg
  vi zoo.cfg
  ```

  修改dataDir, dataLogDir

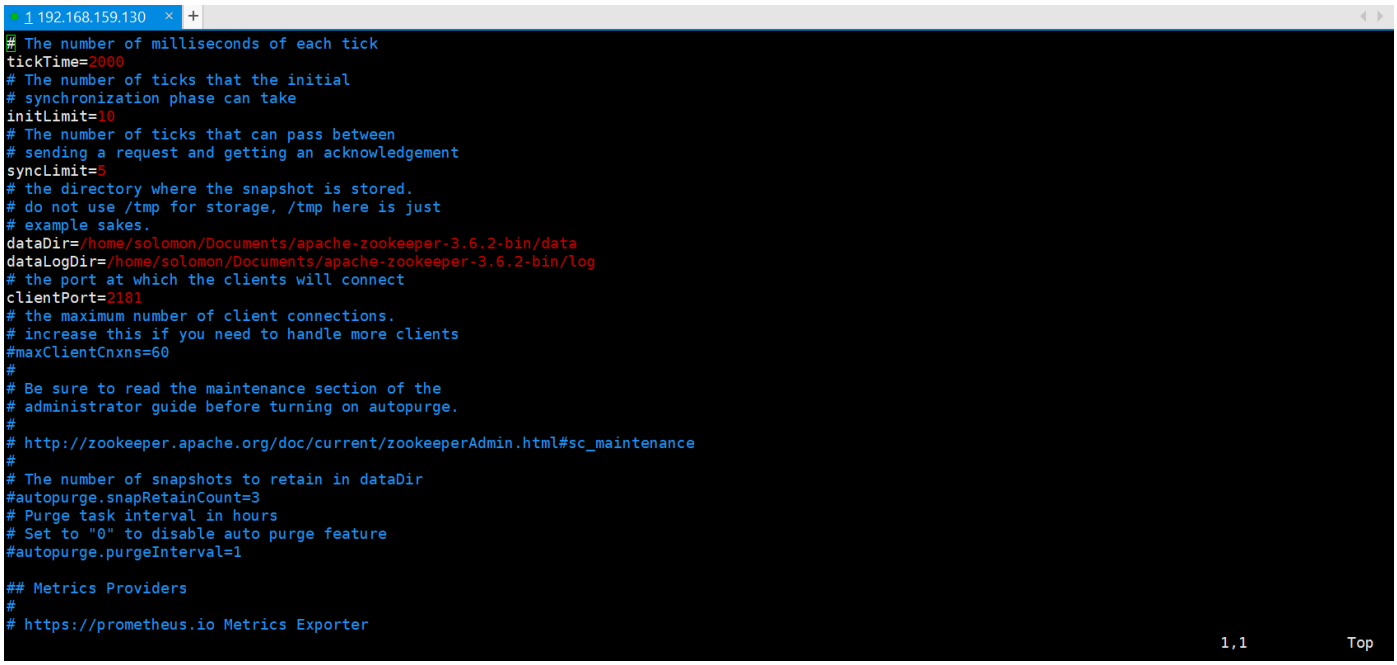

启动命令

```
./zkServer.sh start
```

检查安装命令

```
./zkServer.sh status
```

运行结果如下

- 安装Kafka

到安装包目录下执行命令

```
tar -xvf kafka_2.13-2.6.0.tgz
```

kafka配置文件

```
cd config
vi server.properties
```

修改如下

```
broker.id=1 #三台机器的id分别为1, 2, 3
listeners=PLAINTEXT://192.168.56.130:9092 #设置为每台机器各自的IP
log.dirs=/tmp/kafka-logs
offsets.topic.replication.factor=3
transaction.state.log.replication.factor=3
transaction.state.log.min.isr=1
zookeeper.connect=localhost:2181
```

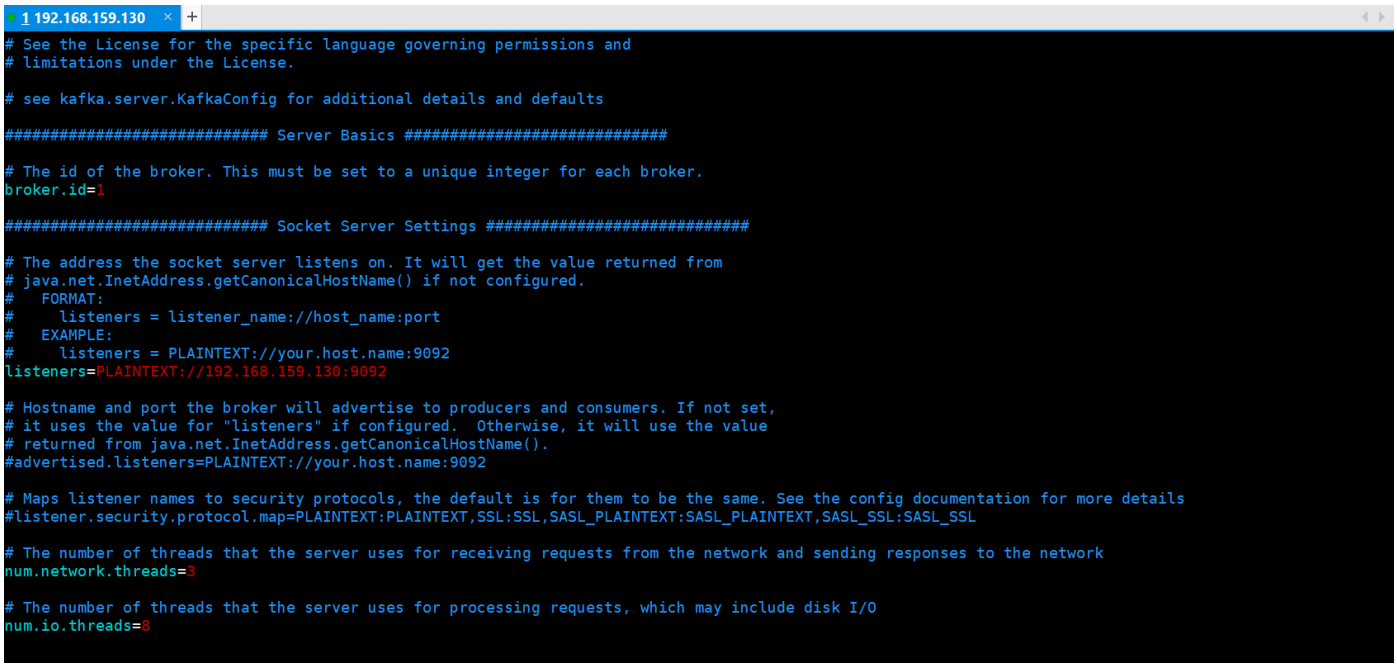

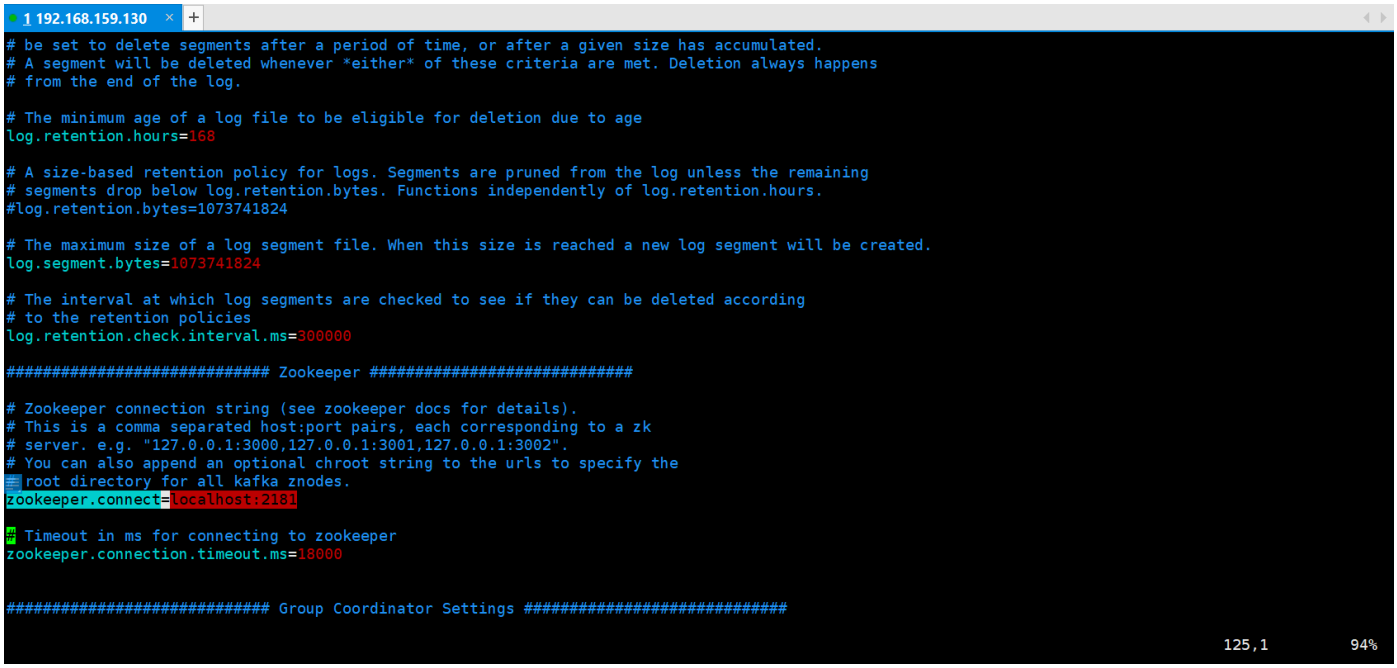

启动

进入安装目录

```
./bin/kafka-server-start.sh -daemon config/server.properties
```

查看状态

```
[3]+  Stopped                 vim server.properties
[solomon@localhost config]$ jps
13447 Jps
3337 QuorumPeerMain
3497 Kafka
[solomon@localhost config]$
```

- Kafka 常用命令工具
  - 创建Topic

    创建一个Topic，使用一个partiton

    ```
    bin/kafka-topics.sh --bootstrap-server 192.168.159.130:9092 --create  --topic myTopic
    ```

    ```
    [solomon@localhost kafka_2.13-2.6.0]$ bin/kafka-topics.sh --bootstrap-server 192.168.159.130:9092 --create  --topic myTopic
    Created topic myTopic.
    [solomon@localhost kafka_2.13-2.6.0]$
    ```

  - 查看创建的Topic

    ```
    bin/kafka-topics.sh --bootstrap-server 192.168.159.130:9092 --describe  --topic myTopic
    ```

    ```
    [solomon@localhost kafka_2.13-2.6.0]$ bin/kafka-topics.sh --bootstrap-server 192.168.159.130:9092 --describe  --topic myTopic
    Topic: myTopic  PartitionCount: 1       ReplicationFactor: 1    Configs: segment.bytes=1073741824
            Topic: myTopic  Partition: 0    Leader: 1       Replicas: 1     Isr: 1
    [solomon@localhost kafka_2.13-2.6.0]$
    ```

  - 生产者

    发送消息进入消息模式

    可以向队列输入消息

    ```
    bin/kafka-console-producer.sh --bootstrap-server 192.168.159.130:9092  --topic myTopic
    ```

    ```
    [solomon@localhost kafka_2.13-2.6.0]$ bin/kafka-console-producer.sh --bootstrap-server 192.168.159.130:9092  --topic myTopic
    >hello pip
    >
    ```

  - 消费者

    接收队列里的消息

    ```
    bin/kafka-console-consumer.sh --bootstrap-server 192.168.159.130:9092  --topic myTopic
    ```

    ```
    [solomon@localhost kafka_2.13-2.6.0]$ bin/kafka-console-consumer.sh --bootstrap-server 192.168.159.130:9092  --topic myTopic

    hello pip
    ```